Breaking the Barriers to Effective Supercomputing: Web Dashboard for Job Accounting and Performance Metrics

Richie Tan
Rosen Center for Advanced Computing
Purdue University
West Lafayette, IN
tan447@purdue.edu

Anjali Rajesh
School of Arts and Sciences
Rutgers University
New Brunswick, NJ
anjali.rajesh@rutgers.edu

Guangzhen Jin
Rosen Center for Advanced Computing
Purdue University
West Lafayette, IN
jin456@purdue.edu

Rajesh Kalyanam
Rosen Center for Advanced Computing
Purdue University
West Lafayette, IN
rkalyana@purdue.edu

***Abstract*—The NSF-funded Anvil supercomputer, built and maintained by the Purdue University Rosen Center for Advanced Computing, enables efficient research computing across a variety of scientific domains nation-wide. In addition to the traditional terminal interface, Anvil uses the open-source web portal framework Open OnDemand to provide a low barrier web interface to the Anvil supercomputer. In this poster we describe enhancements that were made to Anvil's Open OnDemand dashboard to provide a clean, well-structured, and extensible interface for researchers to visualize useful information about their utilization of Anvil without accessing the terminal. Our enhancements include various apps that provide detailed statistics about user jobs and their respective performance metrics while also focusing on data query performance. The information in these apps will enable users to identify and debug any jobs that are noticeably resource-inefficient, as well as improve the queue wait time and efficiency of their jobs in the future.**

***Index Terms*—HPC, Open OnDemand, Anvil, Ruby on Rails, Slurm, dashboard**

Introduction

With the increasing scale of modern data-driven workflows, researchers require efficient yet easy-to-use HPC resources to effectively conduct their research. Many HPC centers typically provide terminal access and utilities for their computing resources but have limited graphical interfaces for researchers with less terminal experience. This project focuses on enhancing the default web-based dashboard for Purdue University's NSF-funded Anvil supercomputer [1], built with Open OnDemand [2]. The goal of the enhanced dashboard is to offer a user-friendly alternative to command-line HPC management, providing an intuitive interface for seamless access to similar data and tools. By developing a user-friendly, productivity-focused dashboard displaying job/allocation accounting data and job performance metrics in an intuitive layout, researchers with varying command line expertise can still easily interact with HPC systems, analyze their workflow performance and resource usage.

System Design

Data Sources

The dashboard retrieves data from two main sources: the Slurm Workload Manager [4] and Purdue Rosen Center for Advanced Computing's (RCAC) Halcyon service [5]. Data from the Slurm database such as job status, accounting information, and partition status is obtained via Slurm shell commands, while Anvil announcements, outages, and maintenance news feeds are retrieved via the Halcyon news API. The disk usage is the only exception, which pulls data from Anvil's ZFS and GPFS servers using a Bash script wrapper called *myquota* (Tab. 1).

Middleware

The middleware for the dashboard consists of a data API and in-memory cache that serves data from the data sources to the Open OnDemand client. Each dashboard feature has a corresponding API endpoint that queries the data sources and stores the data in a memory cache with a certain expiration period. The caching layer ensures that users see as close to real-time data as possible without excessively querying the Slurm or Halcyon databases (Fig. 1).

Dashboard Features

The new dashboard homepage features a welcoming and user-friendly design with widgets and tables that display key metrics: Service Unit (SU) Usage, Disk Usage, the user's current Job Queue, and the status of Anvil's partitions (Fig. 2). Users can monitor resource usage, export data for further analysis, and receive alerts when quotas are reached. The

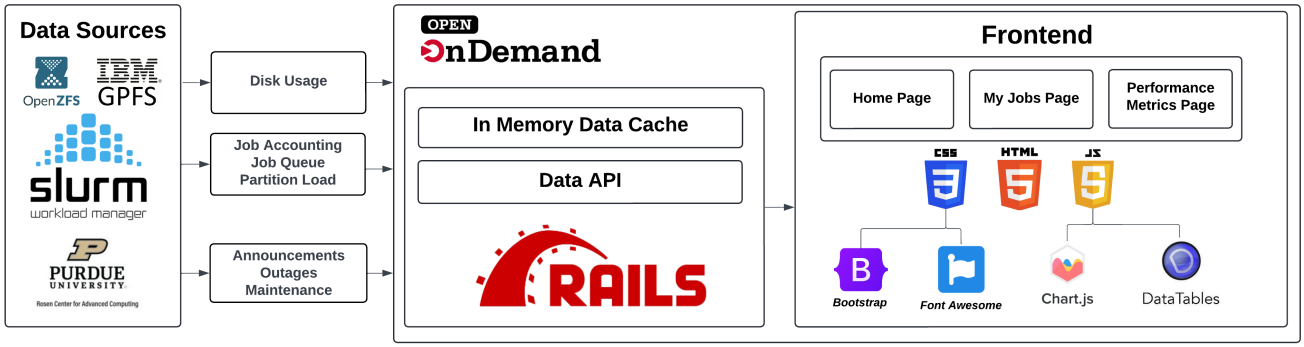

Fig. 1: Flowchart diagram of how the dashboard technologies work together

homepage also provides updates on outages or maintenance that may impact the user.

The *My Jobs* page, accessible from the homepage, displays a table of both the user's own jobs and jobs submitted by other users on the same allocation, including details such as wait time, job efficiency metrics and alerts, real-time SU usage, and a user-friendly job status (e.g. the actual reason a job is pending). Various filtering options are provided in addition to the ability to expand to reveal more information for a job. The page also features graphs depicting SU Usage and Job State distribution for both personal and allocation jobs (Fig. 3).

The *Performance Metrics* page, accessible from the *My Jobs* page, gathers information from the Slurm database to present various job statistics. Users can adjust the data range using tabs for all-time, one year, one month, one week, or one day. The page shows metrics such as the number of jobs within the selected timeframe, the average memory, running time, and CPU efficiencies. The least efficient job ID for each metric is provided so users can further analyze the job to optimize future resource usage (Fig. 4).

Performance Optimizations

The naive solution to showing Slurm job accounting data on the dashboard is to query the Slurm shell commands every time the page is loaded. However, commands like *sacct* and *squeue* fetch data from slurmdbd (Slurm database daemon) and slurmctld (Slurm management daemon) which can become unresponsive under high usage, leading to lags on each web page request. In order to alleviate this, we implemented an in-memory cache that helps to reduce the number of Slurm queries by storing data that doesn't change frequently, including information about completed jobs, disk usage, and service unit usage, while discarding stale data after an expiration period. The My Jobs page first loads the latest cached jobs into the job table before refreshing the data from *sacct*, essentially showing the user nearly up-to-date job data while new data is being fetched while the page finishes loading.

TABLE I: List of Data Sources and Query Commands

Feature	Source	Query
News Feed	Halcyon	GET https://rcac.purdue.edu/api/news
Service Units	Slurm	`scontrol show assoc flags=assoc -o`
Disk Usage	Slurm	`myquota`
Partition Status	Slurm	`sinfo -h -o '%R|%a|%F|%C'`
GPU Count	Slurm	`scontrol show node -oneliner | grep -E 'NodeName=g0'| grep -Eo 'CfgTRES=.* AllocTRES=.* CapWatts'`
Job Queue	Slurm	`squeue -t all -h -o '%i|%P|%j|%u|%T|%r'`
Allocations	Slurm	`sacctmgr show user {username} withassoc format=account -P -n -r | xargs | tr ' ','| tr -d '\n'`
My Jobs	Slurm	`sacct -S now-7days -E now -P -n -a -o {comma separated list of fields`[1]`} -A {comma separated list of user's allocations}`
Performance Metrics	Slurm	`sacct -u {username} -S now-100000days`[2] `-X -o JobID,Submit,Start,End`

Discussion

Through further developments and optimizations, the dashboard can serve as a user-friendly alternative to command-line HPC usage. The dashboard is designed to be adaptable and scalable, allowing for easy configuration for use with other clusters. The data API and caching solution can be used by any HPC resource that utilizes Slurm as its job scheduler or easily adapted to other schedulers such as PBS. HPC centers

[1]jobid,jobname,user,account,start,end,elapsed,submit,planned,partition, state,reason,timelimit,reqmem,alloccpus,totalcpu,workdir,maxrss, maxdiskwrite,maxdiskread,nodelist,reqtres

[2]This represents jobs data from all time.

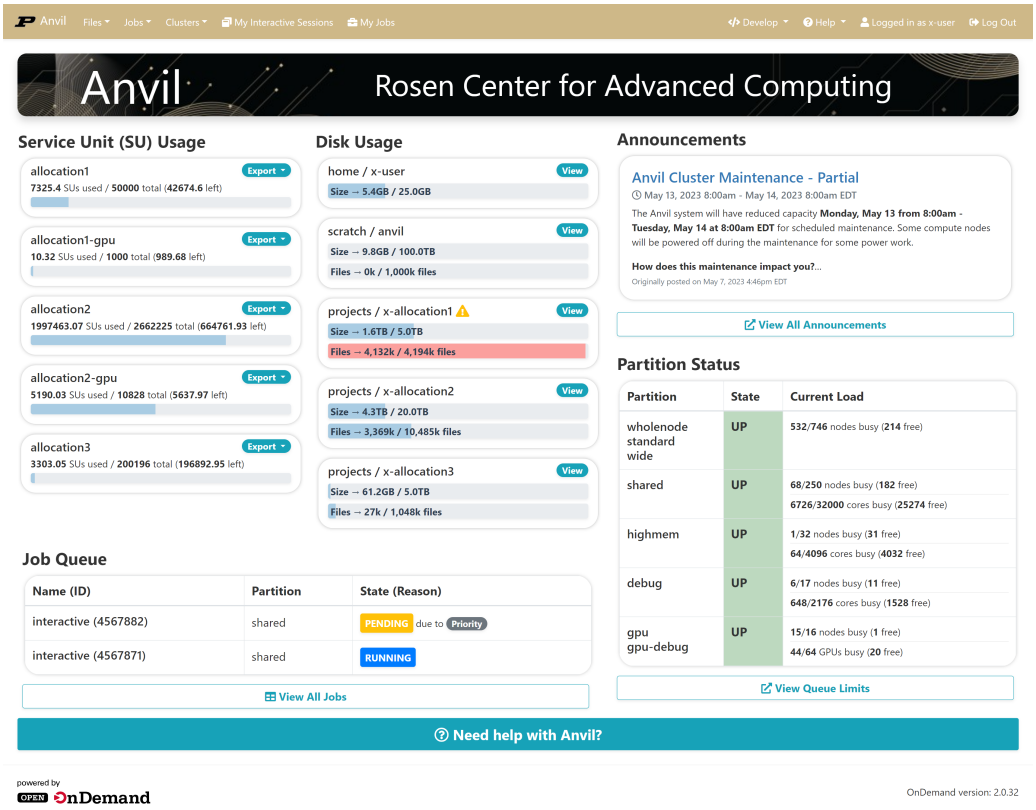

Fig. 2: Anvil dashboard homepage with widgets

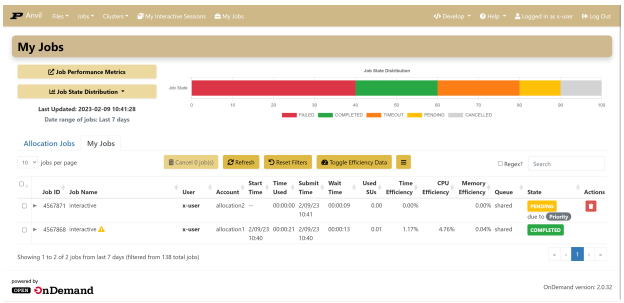

Fig. 3: *My Jobs* app with table of user's jobs

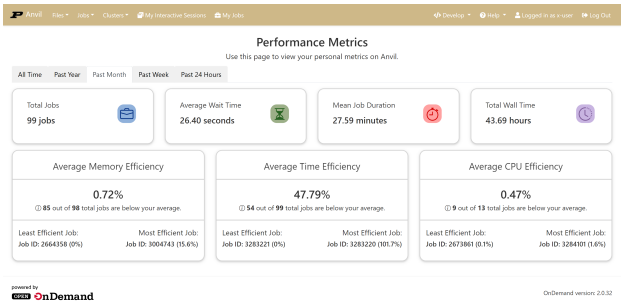

Fig. 4: *Performance Metrics* app with job efficiency and wait time info

that use Open OnDemand can adapt our dashboard framework with modifications to the data sources as needed.

Acknowledgements

This work was supported in part by the National Science Foundation (NSF) grant no. 2005632 and the Anvil Research Experience for Undergraduates (REU) program.

References

[1] X. Carol Song, Preston Smith, Rajesh Kalyanam, Xiao Zhu, Eric Adams, Kevin Colby, Patrick Finnegan, Erik Gough, Elizabett Hillery, Rick Irvine, Amiya Maji, and Jason St. John. 2022. Anvil - System Architecture and Experiences from Deployment and Early User Operations. In Practice and Experience in Advanced Research Computing (PEARC '22). Association for Computing Machinery, New York, NY, USA, Article 23, 1–9. https://doi.org/10.1145/3491418.3530766

[2] Hudak et al., (2018). Open OnDemand: A web-based client portal for HPC centers. Journal of Open Source Software, 3(25), 622, https://doi.org/10.21105/joss.00622

[3] Jette, M.A., Wickberg, T. (2023). Architecture of the Slurm Workload Manager. In: Klusáček, D., Corbalán, J., Rodrigo, G.P. (eds) Job Scheduling Strategies for Parallel Processing. JSSPP 2023. Lecture Notes in Computer Science, vol 14283. Springer, Cham. https://doi.org/10.1007/978-3-031-43943-81

[4] K. D. Colby, D. T. Dietz, P. M. Smith and D. D. Cumberland, "Self-Service Queue and User Management in Shared Clusters," 2014 First International Workshop on HPC User Support Tools, New Orleans, LA, USA, 2014, pp. 22-31, doi: 10.1109/HUST.2014.5. keywords: Delays;Manuals;Engineering profession;Portals;Access control;Business;Conferences,

[5] Rice, S., Colby, K., (2022) Halcyon: Unified HPC Center Operations: An extensible web framework for resource allocation, documentation, and more., Practice and Experience in Advanced Research Computing (PEARC 2022).